

Logistic Regression


Machine Learning
(PM chap 2, chap 3)

Classification

- From $x \rightarrow$ predict y (class)

x

Size of tumor is 0.8 cm



y

cancer ?

Yes / No



Classification

■ Predicting class y

- From x (input value) → predict y (class/category, discrete value)
- *Logistic Regression*: constructing a model $y = f(x)$, using training data




x	class label y
size	cancer
0.75	1 (yes)
0.82	1 (yes)
0.26	0 (no)
0.54	0 (no)
...	...

0.43 → cancer ?



Classification

x		class label
		y
size	weight	cancer
0.75	0.9	1 (yes)
0.82	0.8	1 (yes)
0.26	0.4	0 (no)
0.54	0.7	0 (no)
...




Logistic Regression

■ Example

- Relationship between Age and signs of coronary heart disease (CD)

Age	CD	Age	CD	Age	CD
22	0	40	0	54	0
23	0	41	1	55	1
24	0	46	0	58	1
27	0	47	0	60	1
28	0	48	0	60	0
30	0	49	1	62	1
30	0	49	0	65	1
32	0	50	1	67	1
33	0	51	0	71	1
35	1	51	1	77	1
38	0	52	0	81	1

- linear regression is not appropriate




Logistic Regression

■ The model

- Build a model that predicts the *probability of y*
- $\text{Logit}(p) = f(x)$ is modeled as a linear function

$$\text{logit}(p) = f_{\theta}(x) = w \cdot x + b$$

↑
parameters
(w, b)



x	y
0.75	1
0.82	1
0.26	0
0.54	0
...	...

- w and b are estimated using the training data $(x^{(i)}, y^{(i)})$
 - w : weight (coefficient)
 - b : bias (intercept) θ : parameters
- Minimize prediction error \rightarrow find w and b
 - Gradient descent

Logistic Regression

- Predicting probability of y

- y is usually coded as 1 (true) or 0 (false)
- Predict prob. of y = 1 $p \text{ [0, 1]} = w x + b \text{ [-}\infty, \infty]$ \rightarrow not appropriate


- Odds ratio and Logit

- The **odds** : the ratio of the probability of success to the probability of failure

$$Odds = \frac{p(y = 1|x)}{p(y = 0|x)} \quad [0, \infty]$$

- Logit : log odds

$$Logit(p) = \ln(Odds) = \ln\left(\frac{p}{1-p}\right) \quad [-\infty, \infty]$$



Logistic Regression

- Logistic function


- Solving for P

$$\ln\left(\frac{p}{1-p}\right) = wx + b$$

➡ $\frac{p}{1-p} = e^{wx+b}$

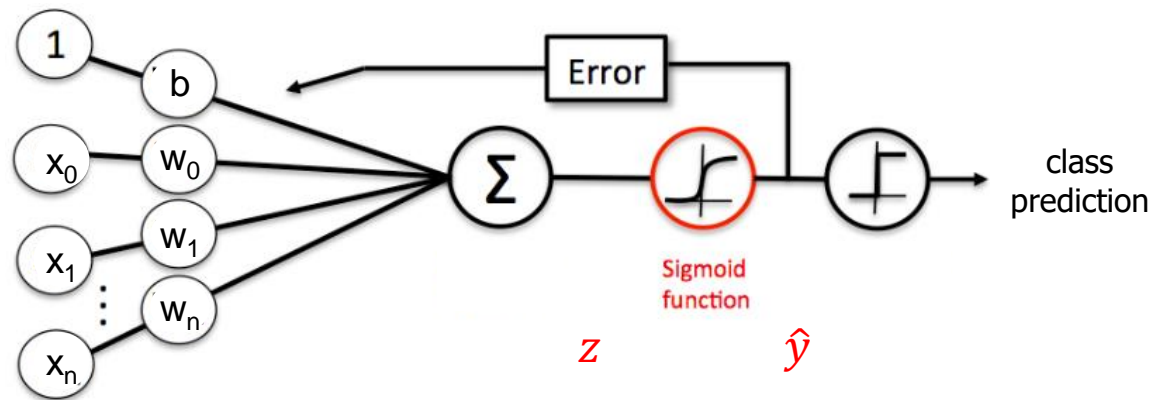
➡ $p = \frac{1}{1 + e^{-(wx+b)}}$

$$= \frac{1}{1 + e^{-z}}$$



Logistic Regression

- The model
 - Output function to predict y : **sigmoid**



$$z = w_0x_0 + w_1x_1 + \dots + b$$

$$\hat{y} = \phi(z) = \frac{1}{1 + e^{-z}}$$

$$\text{class} = \begin{cases} 1 & \text{if } \hat{y} \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Logistic Regression


■ Example

$$\mathbf{x} = (x_0, x_1) \quad \mathbf{w} = (0.1, 0.2) \quad b = -1.4$$

■ Decision boundary

$$\text{class} = 1 \quad \text{if} \quad \hat{y} = \frac{1}{1 + e^{-(0.1x_1 + 0.2x_2 - 1.4)}} \geq 0.5$$

$$\text{class} = 1 \quad \text{if} \quad z = (0.1x_1 + 0.2x_2 - 1.4) \geq 0$$



Estimating Parameters

■ Maximum Likelihood

- For observed data D , model parameter θ , the likelihood $L(\theta) = p(D | \theta)$
- Maximum likelihood estimation (MLE) :

$$\hat{\theta} = \arg \max_{\theta} L(\theta) = \arg \max_{\theta} p(D | \theta)$$

- For i.i.d (independent and identical distributed) data $D = \{d^{(1)}, d^{(2)}, \dots\}$

$$\hat{\theta} = \arg \max_{\theta} p(D | \theta) = \arg \max_{\theta} \prod p(d^{(i)} | \theta)$$

Estimating Parameters

- Maximum likelihood for logistic regression

- The model : $\theta = (w_0, w_1, b)$

$$z = w_0 x_0 + w_1 x_1 + b \quad p = \hat{y} = \frac{1}{1 + e^{-z}} \quad (P(y = 1|x))$$

- For n sample data

$$p(D | \theta) = \prod (p^{(i)})^{y^{(i)}} (1 - p^{(i)})^{1-y^{(i)}} = \prod (\hat{y}^{(i)})^{y^{(i)}} (1 - \hat{y}^{(i)})^{1-y^{(i)}}$$

$$\text{Ex> } \mathbf{y} = (1, 0, 1, 1, 0) \rightarrow p \cdot (1-p) \cdot p \cdot p \cdot (1-p)$$

- Maximizing Likelihood = Minimizing $[-\log(\text{Likelihood})]$

$$\therefore \arg \max_{\theta} p(D | \theta)$$

$$= \arg \min_{\theta} \sum [-y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

Cost Function


■ Data $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$

■ Prediction model $\hat{y} = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + \dots + b)}}$

■ Cost function (binary cross entropy)

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

■ If $y = 1 \quad \Rightarrow \quad -\log(\hat{y})$
 $y = 0 \quad \Rightarrow \quad -\log(1 - \hat{y})$



Gradient of BCE

- For 1 example data

$$J(\theta) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

$$\hat{y} = \frac{1}{1 + e^{-z}}$$

$$z = w_0 x_0 + w_1 x_1 + \dots + b$$

- Gradient of binary cross entropy

$$\frac{\partial J(\theta)}{\partial w_j} = \frac{\partial J(\theta)}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z} \cdot \frac{\partial z}{\partial w_j}$$

$$= \left(-y \frac{1}{\hat{y}} + (1 - y) \frac{1}{(1 - \hat{y})} \right) \cdot \hat{y}(1 - \hat{y}) \cdot x_j$$

$$= (\hat{y} - y) \cdot x_j$$

$$\begin{aligned} \frac{\partial \hat{y}}{\partial z} &= \frac{-(-e^{-z})}{(1 + e^{-z})^2} = \frac{(1 + e^{-z}) - 1}{(1 + e^{-z})^2} \\ &= \frac{1}{(1 + e^{-z})} \cdot \left(1 - \frac{1}{(1 + e^{-z})} \right) \\ &= \hat{y}(1 - \hat{y}) \end{aligned}$$

GD for Logistic Regression

- Data $(x_0^{(1)}, x_1^{(1)}, \dots, y^{(1)}), (x_0^{(2)}, x_1^{(2)}, \dots, y^{(2)}), \dots, (x_0^{(m)}, x_1^{(m)}, \dots, y^{(m)})$

- Model $\hat{y} = \frac{1}{1 + e^{-z}} \quad z = w_0 x_0 + w_1 x_1 + \dots + b$

- Cost $J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$

- Gradients

$$\frac{\partial J}{\partial w_0} = \frac{1}{m} \sum (\hat{y}^{(i)} - y^{(i)}) x_0^{(i)}$$

$$\frac{\partial J}{\partial w_1} = \frac{1}{m} \sum (\hat{y}^{(i)} - y^{(i)}) x_1^{(i)}$$

...

$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum (\hat{y}^{(i)} - y^{(i)})$$

GD for Logistic Regression

- Vector form - Data (N features, M data)

$$\mathbf{X} = \begin{matrix} & \begin{bmatrix} x_0^{(1)} & x_1^{(1)} & \dots \\ x_0^{(2)} & x_1^{(2)} & \dots \\ x_0^{(3)} & x_1^{(3)} & \dots \\ \dots & \dots & \dots \end{bmatrix} \\ \text{(M x N)} & \end{matrix} \quad \mathbf{y} = \begin{matrix} & \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(3)} \\ \dots \end{bmatrix} \\ \text{(M x 1)} & \end{matrix} \quad \mathbf{w} = \begin{matrix} & \begin{bmatrix} w_0 \\ w_1 \\ \dots \end{bmatrix} \\ \text{(N x 1)} & \end{matrix}$$

- Model

$$z^{(1)} = w_0 x_0^{(1)} + w_1 x_1^{(1)} \dots + b$$

$$z^{(2)} = w_0 x_0^{(2)} + w_1 x_1^{(2)} \dots + b$$

$$z^{(3)} = w_0 x_0^{(3)} + w_1 x_1^{(3)} \dots + b$$

$$\dots$$
$$\hat{y}^{(i)} = \text{sigmoid}(z^{(i)})$$

- Cost

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

$$\begin{bmatrix} z^{(1)} \\ z^{(2)} \\ z^{(3)} \\ \dots \end{bmatrix} = \begin{bmatrix} x_0^{(1)} & x_1^{(1)}, \dots \\ x_0^{(2)} & x_1^{(2)}, \dots \\ x_0^{(3)} & x_1^{(3)}, \dots \\ \dots & \dots \end{bmatrix} \cdot \begin{bmatrix} w_0 \\ w_1 \\ \dots \end{bmatrix} + b$$

$$\mathbf{z} = \mathbf{X} \cdot \mathbf{w} + b$$

$$\hat{\mathbf{y}} = \text{sigmoid}(\mathbf{z})$$

$$\text{cost} = \frac{1}{m} \text{sum}(-\mathbf{y} \log(\hat{\mathbf{y}}) - (1 - \mathbf{y}) \log(1 - \hat{\mathbf{y}}))$$

GD for Logistic Regression

■ Gradient

$$\frac{\partial J}{\partial w_0} = \frac{1}{m} \left((\hat{y}^{(1)} - y^{(1)})x_0^{(1)} + (\hat{y}^{(2)} - y^{(2)})x_0^{(2)} + \dots \right)$$

$$\frac{\partial J}{\partial w_1} = \frac{1}{m} \left((\hat{y}^{(1)} - y^{(1)})x_1^{(1)} + (\hat{y}^{(2)} - y^{(2)})x_1^{(2)} + \dots \right)$$

...

$$\frac{\partial J}{\partial b} = \frac{1}{m} \left((\hat{y}^{(1)} - y^{(1)}) + (\hat{y}^{(2)} - y^{(2)}) + \dots \right)$$

$$\begin{bmatrix} \frac{\partial J}{\partial w_0} \\ \frac{\partial J}{\partial w_1} \\ \vdots \end{bmatrix} = \frac{1}{m} \begin{bmatrix} x_0^{(1)}, x_0^{(2)}, x_0^{(3)}, \dots \\ x_1^{(1)}, x_1^{(2)}, x_1^{(3)}, \dots \\ \dots \end{bmatrix} \cdot \begin{bmatrix} \hat{y}^{(1)} - y^{(1)} \\ \hat{y}^{(2)} - y^{(2)} \\ \hat{y}^{(3)} - y^{(3)} \\ \dots \end{bmatrix}$$

$$\nabla J(\mathbf{w}) = \frac{1}{m} \mathbf{X}^T (\hat{\mathbf{y}} - \mathbf{y})$$

$$\frac{\partial J}{\partial b} = \frac{1}{m} [1 \quad 1 \quad 1] \cdot \begin{bmatrix} \hat{y}^{(1)} - y^{(1)} \\ \hat{y}^{(2)} - y^{(2)} \\ \hat{y}^{(3)} - y^{(3)} \\ \dots \end{bmatrix}$$

$$\frac{\partial J}{\partial b} = \frac{1}{m} \text{sum}(\hat{\mathbf{y}} - \mathbf{y})$$

GD for Logistic Regression

- Learning (gradient descent)

Given

$$\mathbf{X} = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} & \dots \\ x_0^{(2)} & x_1^{(2)} & \dots \\ x_0^{(3)} & x_1^{(3)} & \dots \\ \dots & \dots & \dots \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(3)} \\ \dots \end{bmatrix}$$

Initialize

$$\mathbf{w} = [w_1 \quad w_2 \quad \dots], \quad b$$

repeat

$$\mathbf{z} = \mathbf{X} \cdot \mathbf{w} + b$$
$$\hat{\mathbf{y}} = \text{sigmoid}(\mathbf{z})$$

*decision
boundary*

$$\mathbf{w} = \mathbf{w} - \alpha \frac{1}{m} (\mathbf{X}^T (\hat{\mathbf{y}} - \mathbf{y}))$$

$$b = b - \alpha \frac{1}{m} (\text{sum}(\hat{\mathbf{y}} - \mathbf{y}))$$

$$\mathbf{w} = \mathbf{w} - \alpha \nabla J(\mathbf{w})$$

$$b = b - \alpha \frac{\partial J}{\partial b}$$

$$\text{cost} = \frac{1}{m} \text{sum}(-\mathbf{y} \log(\hat{\mathbf{y}}) - (1 - \mathbf{y}) \log(1 - \hat{\mathbf{y}}))$$

Example

■ Data


```
X_train = np.array([[0.5, 1.5], [1.0, 1.0], [1.0, 0.5],  
                    [2.5, 1.5], [2.0, 2.0], [1.0, 2.5]])  
y_train = np.array([0, 0, 0, 1, 1, 1])
```

■ Compute \hat{y} and cost

```
z = np.dot(X, w) + b  
y_hat = sigmoid(z)  
  
cost = -y * np.log(y_hat) - (1 - y) * np.log(1 - y_hat)  
cost = np.sum(cost) / m
```

■ Compute gradients

```
y_hat = sigmoid(np.dot(X, w) + b)  
err = y_hat - y  
  
dj_dw = np.dot(X.T, err) / m  
dj_db = np.sum(err) / m
```



$$\mathbf{z} = \mathbf{X} \cdot \mathbf{w} + b$$

$$\hat{y} = \text{sigmoid}(\mathbf{z})$$

$$\text{cost} = \frac{1}{m} \sum (-y \log(\hat{y}) - (1 - y) \log(1 - \hat{y}))$$

$$\frac{\partial J}{\partial \mathbf{w}} = \frac{1}{m} \mathbf{X}^T (\hat{\mathbf{y}} - \mathbf{y})$$

$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum (\hat{y} - y)$$

Example

- Learning (gradient descent)

```
def gradient_descent(X, y, w, b, alpha, num_iters):  
    J_history = []  
    for i in range(num_iters):  
        dj_db, dj_dw = compute_gradient(X, y, w, b)  
  
        w = w - alpha * dj_dw  
        b = b - alpha * dj_db  
  
        J_history.append(compute_cost(X, y, w, b) )
```

$$\mathbf{w} = \mathbf{w} - \alpha \nabla J(\mathbf{w})$$
$$b = b - \alpha \frac{\partial J}{\partial b}$$

```
w_init = np.zeros(X_train.shape[1])  
b_init = 0.  
alpha = 0.1  
iterations = 10000  
  
w_final, b_final, J_hist = gradient_descent(X_train, y_train, w_init, b_init, alpha, iterations)
```


```
Iteration 0: Cost 0.6813565790579484  
Iteration 1000: Cost 0.12340171021257584  
Iteration 2000: Cost 0.06626538139036627  
Iteration 3000: Cost 0.045349110494241084  
Iteration 4000: Cost 0.034514148702862464  
Iteration 5000: Cost 0.02788411906632822
```

$\mathbf{w} = [4.3876931 \quad 5.03096199], \quad b = -13.24166197110235$

Example


- Cost change

```
plt.plot(J_hist[:10000])
```



- Learned model


```
fig, ax = plt.subplots(subplot_kw={"projection": "3d"})  
  
ax.scatter(X_train[:3,0], X_train[:3,1], y_train[:3], marker='o', c='r')  
ax.scatter(X_train[3:,0], X_train[3:,1], y_train[3:], marker='x', c='b')  
  
x0 = np.arange(0, 3, 0.1)  
x1 = np.arange(0, 3, 0.1)  
x0, x1 = np.meshgrid(x0, x1)  
  
y_hat = sigmoid(x0 * w_final[0] + x1 * w_final[1] + b_final)  
ax.plot_surface(x0, x1, y_hat, cmap=cm.coolwarm, alpha=0.5)
```



Example

- Learned decision boundary

```
x0 = np.arange(0,4)
x1 = -w_final[0] / w_final[1] * x0 - b_final / w_final[1]
plt.plot(x0, x1)
```



$$\mathbf{z} = \mathbf{x} \cdot \mathbf{w} + b = 0$$
$$w_0x_0 + w_1x_1 + b = 0$$

- Accuracy of the model

```
y_pred = predict(X_train, w_final, b_final)
accuracy = np.sum(y_train == y_pred) / len(y_train)
```

Example

- Preventing overflow
 - Prevent $\exp(\infty)$ in sigmoid or softmax

```
# sigmoid function  
def sigmoid(z):  
    1. / (1. + np.exp(-np.clip(z, -250, 250))) # np.clip to prevent overflow
```

- Prevent $\log(0)$ in computing cross entropy

```
# cross entropy  
def compute_cost(y, o):  
    np.sum(-y*(np.log(o+1e-7))) # +1e-7 to prevent overflow
```

Logistic Regression using Scikit Learn

- Load iris dataset from scikit learn → return dictionary with attributes:


```
from sklearn import datasets  
iris = datasets.load_iris()
```

- data
- feature_names
- target
- target_names

```
[6.5 3. 5.2 2. ]  
[6.2 3.4 5.4 2.3]  
[5.9 3. 5.1 1.8]]  
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']  
[[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
 2 2]  
['setosa' 'versicolor' 'virginica']
```

- Get X, y data
 - Select only features 2, 3
 - Select only class 0 and 1

```
X = iris.data[0:100, [2, 3]]
y = iris.target[0:100]
```



Logistic Regression using Scikit Learn

- Load iris dataset from scikit learn → return dictionary with attributes:

```
from sklearn import datasets  
iris = datasets.load_iris()
```

- data
- feature_names
- target
- target_names

```
[6.5 3.  5.2 2. ]  
[6.2 3.4 5.4 2.3]  
[5.9 3.  5.1 1.8]]  
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']  
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
 2 2]  
['setosa' 'versicolor' 'virginica']
```


Logistic Regression using Scikit Learn

- Read into a DataFrame

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
...
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

- Get X, y array
 - Select only features 2, 3
 - Select only class 0 and 1

```
X = iris.data[0:100, [2, 3]]  
y = iris.target[0:100]
```



Logistic Regression using Scikit Learn

- Split data into training / test dataset
 - sklearn.model_selection.train_test_split

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size=0.3, random_state=1, stratify=y)  
print(X_train.shape)  
print(X_test.shape)
```

```
(70, 2)  
(30, 2)
```

- Use [LogisticRegression](#). Training the model - `.fit()`


```
lr = LogisticRegression(C=100.0)  
lr.fit(X_train, y_train)  
  
print('w = ', lr.coef_)  
print('b = ', lr.intercept_)
```

```
w = [[5.52734478 2.26785607]]  
b = [-16.33236067]
```

Logistic Regression using Scikit Learn

- Plotting the decision boundary
 - Use np.meshgrid, plt.contour

```
xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),  
                        np.arange(x2_min, x2_max, resolution))  
  
Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)  
Z = Z.reshape(xx1.shape)  
plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)  
  
for idx, c1 in enumerate(np.unique(y)):  
    plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1],  
                alpha=0.8, c=colors[idx], marker=markers[idx], label=c1)
```



Logistic Regression using Scikit Learn

- Accuracy of the learned model - `.score()`

```
print('Training accuracy: %.2f' % lr.score(X_train, y_train))  
print('Test accuracy: %.2f' % lr.score(X_test, y_test))
```

Training accuracy: 1.00
Test accuracy: 1.00

- Computing class probability - `.predict_proba()`

```
print(lr.predict_proba(X_test[:5]))
```

```
[[2.70296563e-06 9.99997297e-01]  
 [6.04085017e-02 9.39591498e-01]  
 [9.99767710e-01 2.32289833e-04]  
 [9.99596350e-01 4.03650310e-04]  
 [9.99767710e-01 2.32289833e-04]]
```

- Predicting class labels - `.predict()`

```
print('True test labels :', y_test[:5])  
print('Predicted labels :', lr.predict(X_test[:5]))
```

True test labels : [1 1 0 0 0]
Predicted labels : [1 1 0 0 0]

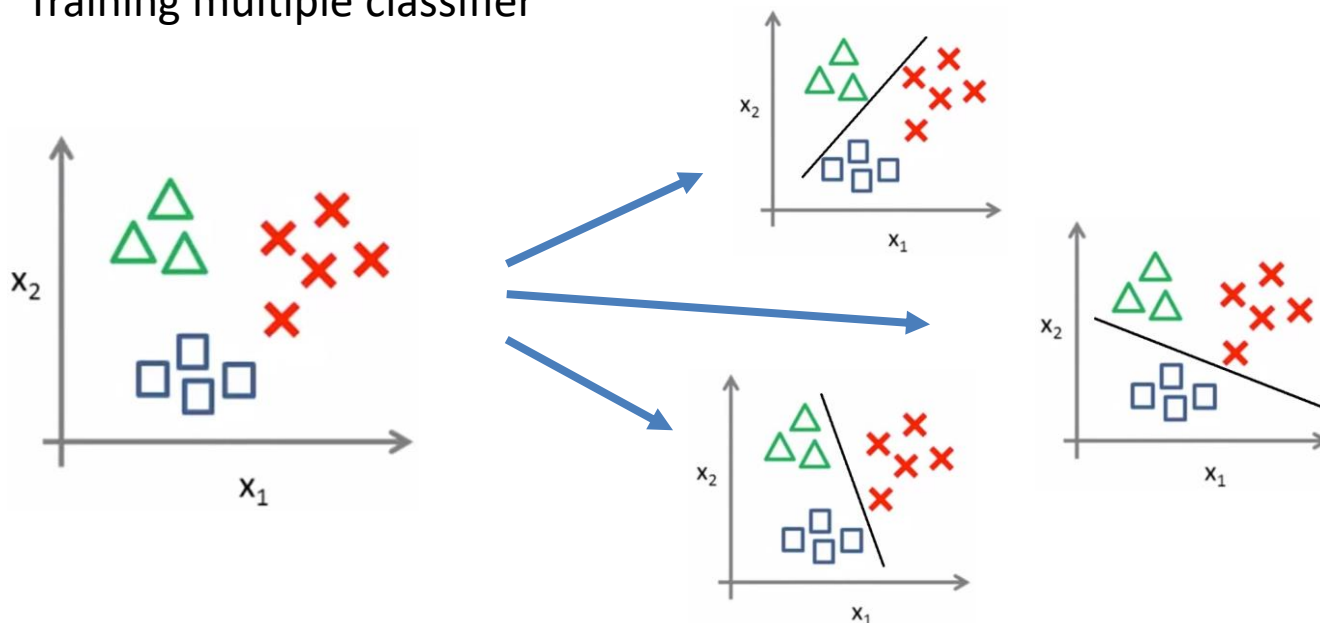
Multinomial Logistic Regression

- Multiple class case

- Classes $\{0, 1\}$: binary classification
- Classes $\{A, B, C\}$: multinomial classification $\rightarrow \{(1, 0, 0), (0, 1, 0), (0, 0, 1)\}$

- OvR (One vs. Rest)


- Training multiple classifier



Multinomial Logistic Regression

- OvR (One vs. Rest)

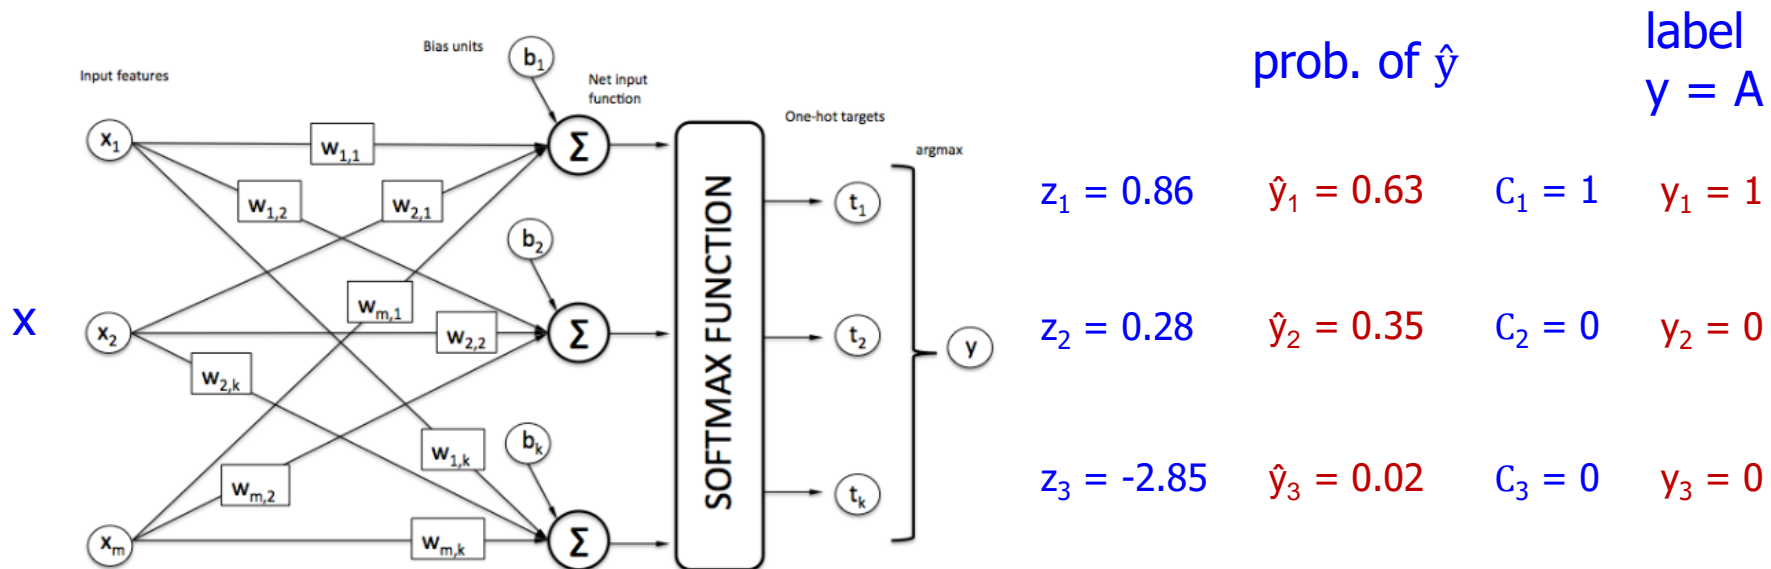
- Decision boundary combining 3 classifier



Multinomial Logistic Regression

■ Multinomial

- Output function to predict y : **softmax**



$$z = w_0x_1 + w_1x_1 + w_2x_2 + \dots + b$$

$$\hat{y}_k = \frac{e^{z_k}}{\sum_i e^{z_i}}$$


$$Class_k = \begin{cases} 1 & \text{if } \hat{y}_k \text{ is max} \\ 0 & \text{otherwise} \end{cases}$$

Cost Function

- Cost function (cross entropy)

$$J(\theta) = \sum J^{(i)}(\theta)$$

$$J^{(i)}(\theta) = - \sum_k y_k \log \hat{y}_k$$



- Example

$$y_1 = 1 \quad \hat{y}_1 = 0.63$$

$$y_2 = 0 \quad \hat{y}_2 = 0.25$$

$$y_3 = 0 \quad \hat{y}_3 = 0.12$$



$$-\sum y \log \hat{y} = - (1 \times \log 0.63 + 0 \times \log 0.25 + 0 \times \log 0.12)$$

※ For binary classification

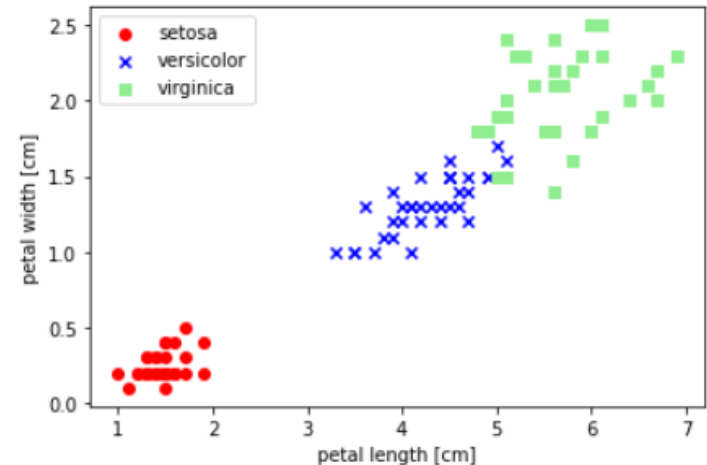
$$y = 1 \quad \hat{y} = 0.63$$



$$-\sum y \log \hat{y} = - (1 \times \log 0.63 + 0 \times \log 0.37)$$

Multinomial Logistic Regression using Scikit Learn


- Load dataset from scikit learn
 - Select only features 2, 3
 - Select all class 0, 1, 2
- Split data into training test dataset
- Training the model



```
lr = LogisticRegression(C=100.0, random_state=1, multi_class='ovr')
lr.fit(X_train, y_train)
```

```
print('w = ', lr.coef_)
print('b = ', lr.intercept_)
```

```
w = [[-5.52741894 -2.26767352]
      [ 1.34492291 -2.71926866]
      [ 6.95808994  7.36736804]]
b = [ 16.33234585 -2.59905851 -46.45251225]
```



Multinomial Logistic Regression using Scikit Learn

- The decision boundary

- Computing class probability

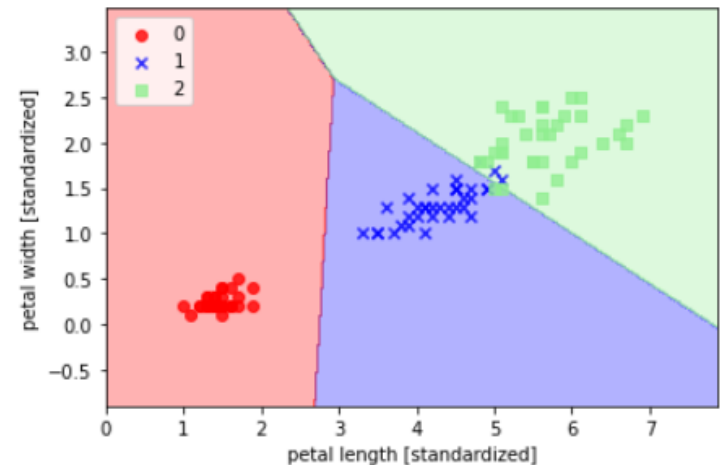
```
print(lr.predict_proba(X_test[:5]))
```

```
[[6.27006745e-09 1.44806162e-01 8.55193832e-01]
 [8.34564702e-01 1.65435298e-01 1.46365975e-14]
 [8.49059341e-01 1.50940659e-01 8.82090981e-16]
 [1.35565353e-05 7.76776434e-01 2.23210009e-01]
 [3.69241060e-05 9.89606981e-01 1.03560948e-02]]
```

- Predicting class labels

```
print('True test labels :', y_test[:5])
print('Predicted labels :', lr.predict(X_test[:5]))
```

```
True test labels : [2 0 0 2 1]
Predicted labels : [2 0 0 1 1]
```



Regularization

- Regularization

- Add regularization term(*penalty*) to a cost function to prevent overfitting

- L2 regularization (ridge)

- Add

$$\lambda \|\mathbf{w}\|_2^2 = \lambda \sum_{j=1}^n w_j^2 = \lambda(w_1^2 + w_2^2 + \dots + w_n^2)$$

- Makes weights small, but not to 0. Less computationally expensive
- Gradient descent

$$w_j = w_j - \alpha \frac{\partial J}{\partial w_j} - \alpha \lambda w_j$$

Regularization

- L1 regularization (lasso)

- Add


$$\lambda \|\mathbf{w}\|_1 = \lambda \sum_{j=1}^n |w_j| = \lambda(|w_1| + |w_2| + \cdots + |w_n|)$$

- Makes some weights to 0 \rightarrow feature selection. Robust to outliers
 - Gradient descent

$$w_j = w_j - \alpha \frac{\partial J}{\partial w_j} - \alpha \lambda \text{sign}(w_j)$$

Regularization

■ Lasso vs Ridge




L2 Regularization in Scikit Learn

- Scikit learn **LogisticRegression** parameter **C**

$$C = \frac{1}{\lambda}$$

$$J(\theta) = C \left[\sum_{i=1}^n [-y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] \right] + \frac{1}{2} \|\mathbf{w}\|_2^2$$

- Effect of C
 - Small C \rightarrow large $\lambda \rightarrow$ small w




L2 Regularization in Scikit Learn


- Training the model with various C values

```
for c in np.arange(-5, 5):  
    lr = LogisticRegression(C=10.**c,  
                           random_state=1,  
                           multi_class='ovr')  
    lr.fit(X_train_std, y_train)  
  
    params.append(10.**c)  
    weights.append(lr.coef_[1])  
    test_acc.append(lr.score(X_test_std, y_test))
```


- Plot the parameter change



lr.coef_[0]



lr.coef_[1]



lr.coef_[2]